



UNIVERSIDADE  
CATÓLICA  
PORTUGUESA

BRAGA

# Deep Learning

Session 5

## Training Neural Networks

Applied Data Science

2024/2025

# Objective Function: Analogous to Training

Babies to not throw food on the floor



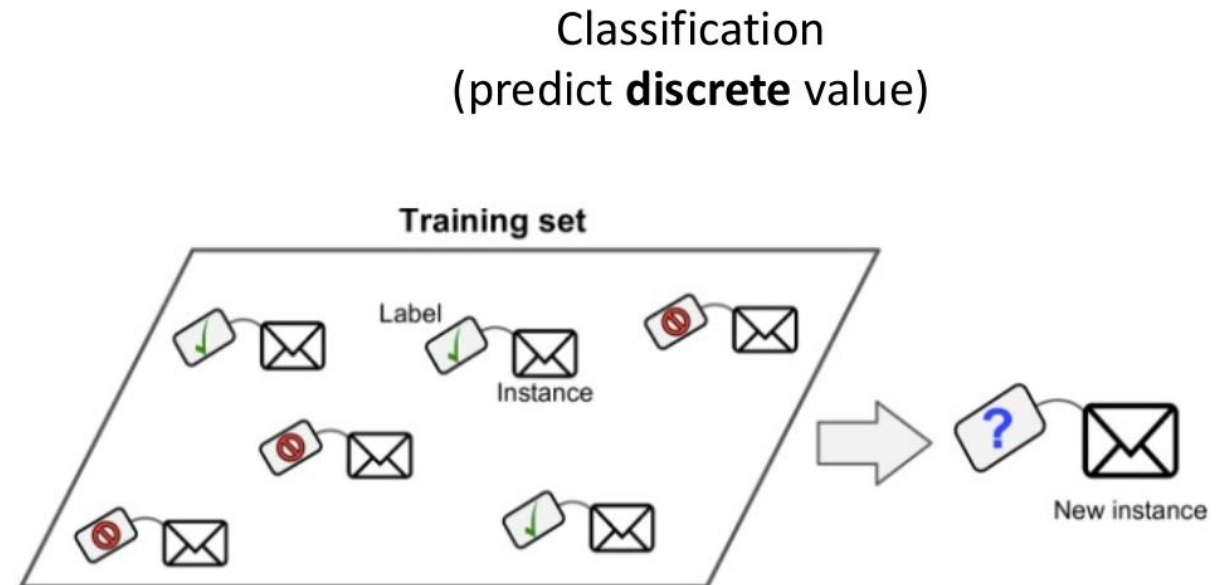

<https://www.youtube.com/watch?v=58Pr-QrVNqU>

Dogs to learn to sit




<https://www.akc.org/expert-advice/training/how-to-become-a-dog-trainer/>

# Objective Function: Learn Model Parameters that Achieve a Specified Goal



# Objective Function: Learn Model Parameters that Achieve a Specified Goal



e.g., make as small as possible the squared error (aka, L2 loss, quadratic loss)

Mean taken over  $n$  instances

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

True value


Predicted value

What is the range of possible values?

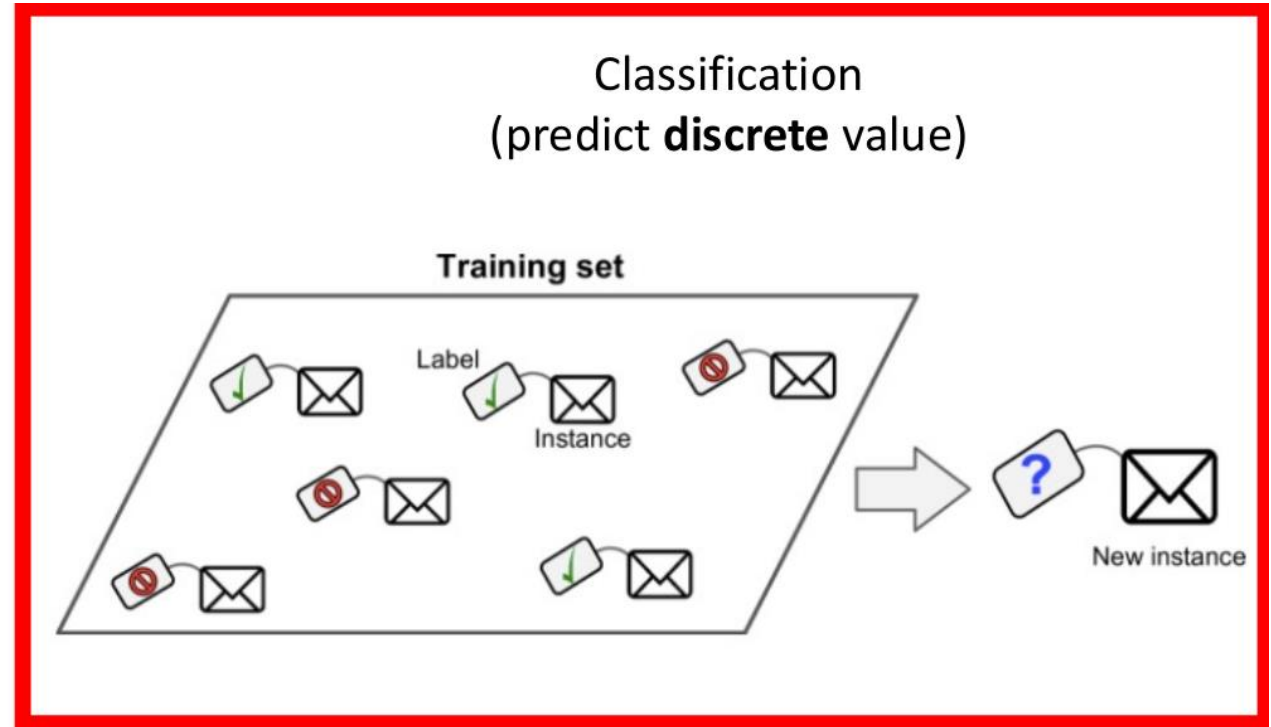
- Minimum: 0
  - i.e., all correct predictions
- Maximum: Infinity
  - i.e., incorrect predictions

# Objective Function: Learn Model Parameters that Achieve a Specified Goal

Regression  
(predict **continuous** value)




Classification  
(predict **discrete** value)





# Objective Function: Learn Model Parameters that Achieve a Specified Goal



e.g., make as small as possible the distance between predicted and true class distributions with **cross entropy loss**

True prediction is 1 hot vector (i.e., one 1 and the rest 0s)

<https://towardsdatascience.com/multi-label-image-classification-with-neural-network-keras-ddc1ab1afede>

# Objective Function: Learn Model Parameters that Achieve a Specified Goal

Probability distribution of predicted class

Probability distribution of true class

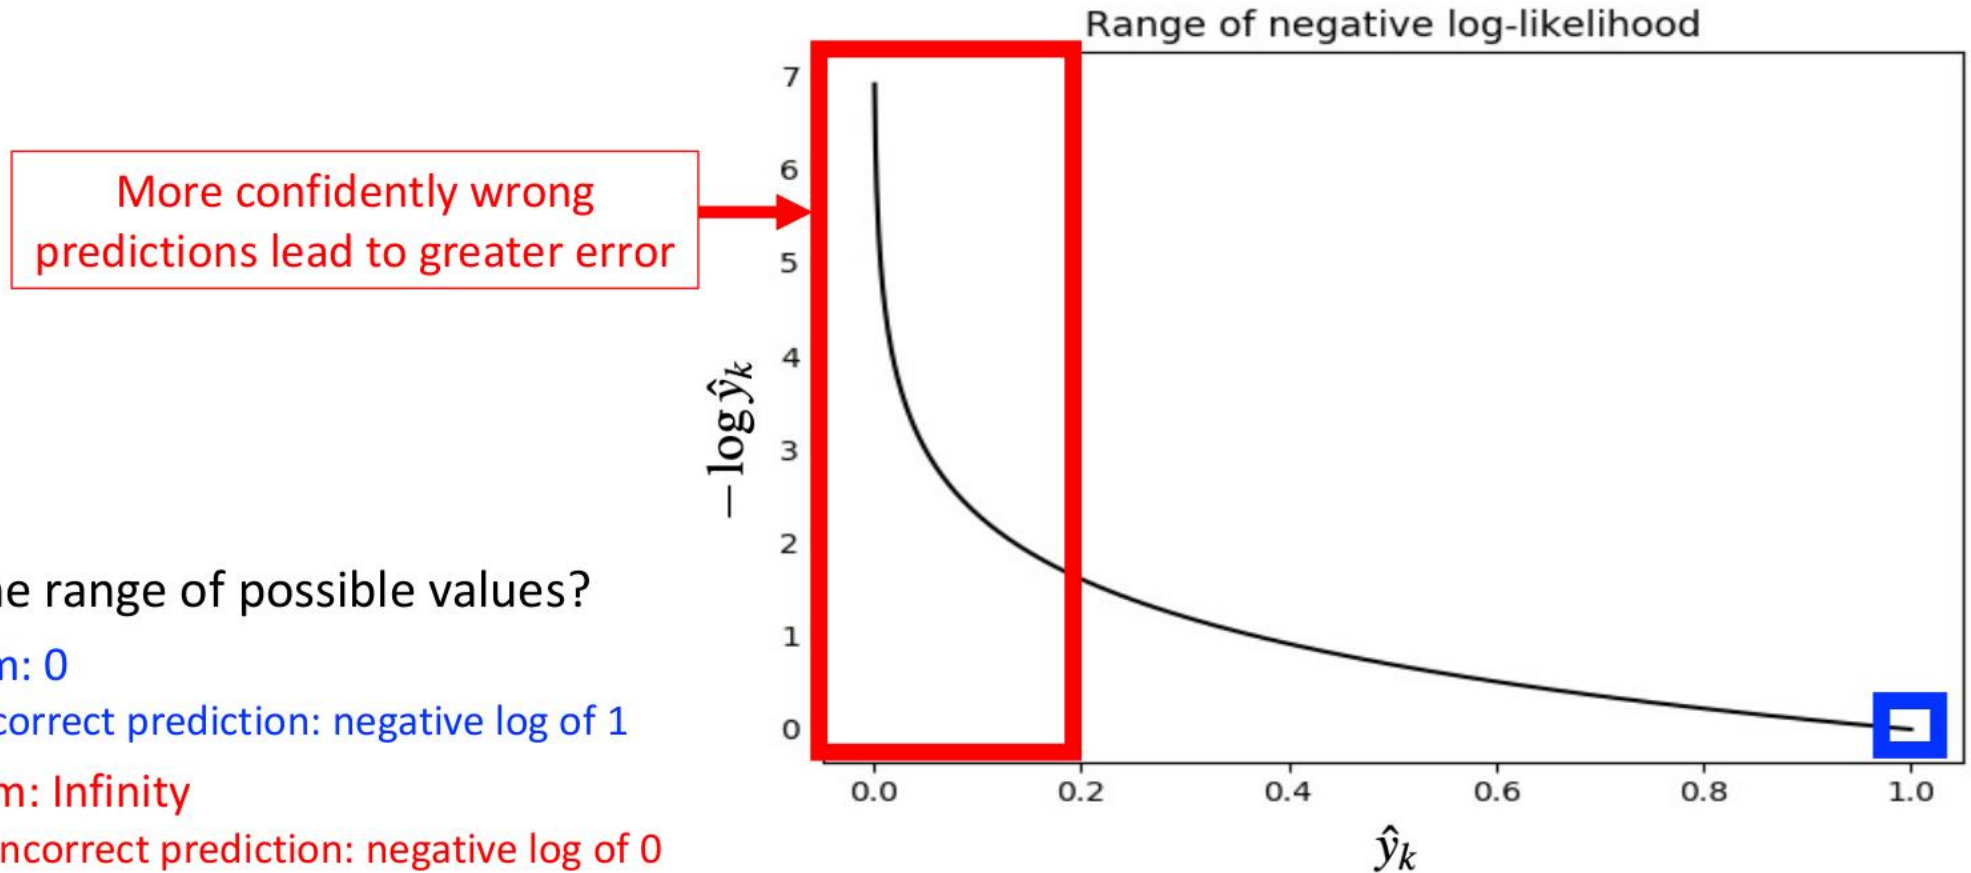
Number of classes?

$$L_{\text{CE}}(\hat{y}, y) = - \sum_{k=1}^K y_k \log \hat{y}_k$$

What is the range of possible values?

- Minimum: 0
  - i.e., correct prediction: negative log of 1
- Maximum: Infinity
  - i.e., incorrect prediction: negative log of 0

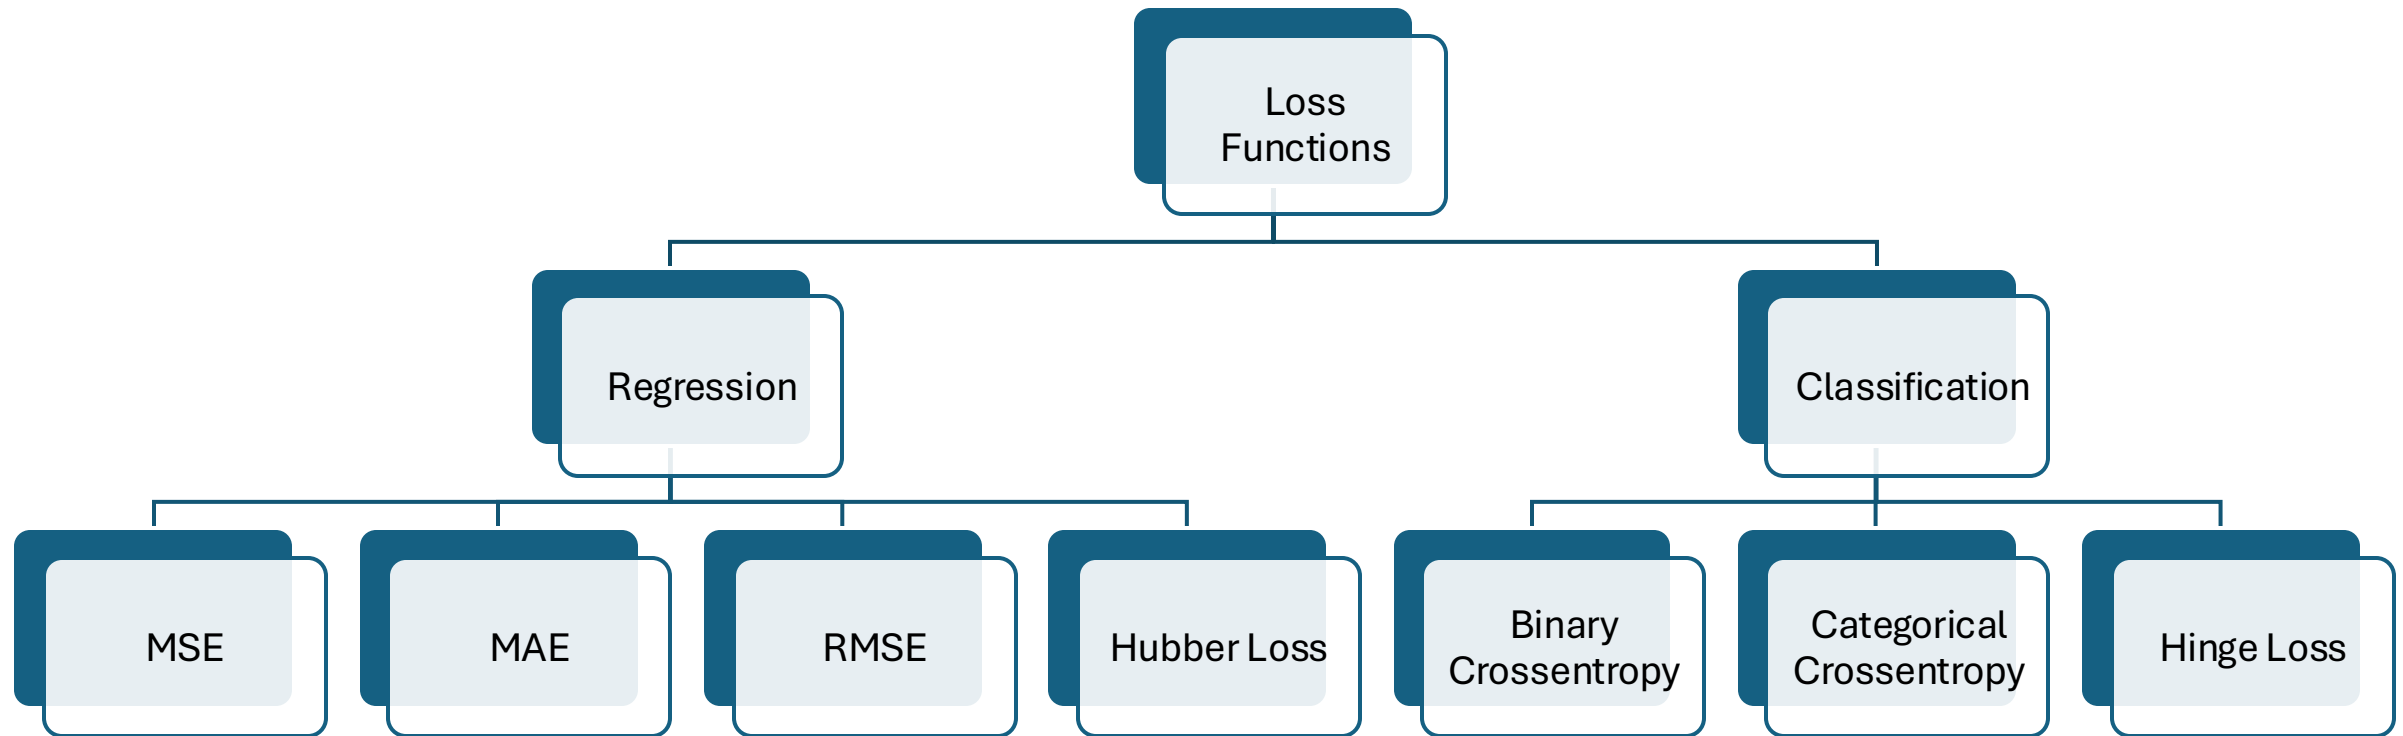
# Objective Function: Learn Model Parameters that Achieve a Specified Goal





# Objective Function

- Many objective functions exist, and we will explore popular ones in this course!



# Loss Functions



## 10 Most Common Loss Functions in Machine Learning

 [blog.DailyDoseofDS.com](https://blog.dailydoseofds.com)


Loss Function Name	Description	Function
<b>Regression Loss Functions</b>		
Mean Bias Error	Captures average bias in prediction. But is rarely used for training.	$\mathcal{L}_{MBE} = \frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))$
Mean Absolute Error	Measures absolute average bias in prediction. Also called L1 Loss.	$\mathcal{L}_{MAE} = \frac{1}{N} \sum_{i=1}^N  y_i - f(x_i) $
Mean Squared Error	Average squared distance between actual and predicted. Also called L2 Loss.	$\mathcal{L}_{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))^2$
Root Mean Squared Error	Square root of MSE. Loss and dependent variable have same units.	$\mathcal{L}_{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))^2}$
Huber Loss	A combination of MSE and MAE. It is parametric loss function.	$\mathcal{L}_{Huberloss} = \begin{cases} \frac{1}{2}(y_i - f(x_i))^2 & :  y_i - f(x_i)  \leq \delta \\ \delta( y_i - f(x_i)  - \frac{1}{2}\delta) & : otherwise \end{cases}$
Log Cosh Loss	Similar to Huber Loss + non-parametric. But computationally expensive.	$\mathcal{L}_{LogCosh} = \frac{1}{N} \sum_{i=1}^N \log(\cosh(f(x_i) - y_i))$
<b>Classification Loss Functions (Binary + Multi-class)</b>		
Binary Cross Entropy (BCE)	Loss function for binary classification tasks.	$\mathcal{L}_{BCE} = \frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(x_i)) + (1 - y_i) \cdot \log(1 - p(x_i))$
Hinge Loss	Penalizes wrong and right (but less confident) predictions. Commonly used in SVMs.	$\mathcal{L}_{Hinge} = \max(0, 1 - (f(x) \cdot y))$
Cross Entropy Loss	Extension of BCE loss to multi-class classification.	$\mathcal{L}_{CE} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \cdot \log(f(x_{ij}))$ <i>N : samples; M : classes</i>
KL Divergence	Minimizes the divergence between predicted and true probability distribution	$\mathcal{L}_{KL} = \sum_{i=1}^N y_i \cdot \log\left(\frac{y_i}{f(x_i)}\right)$

<https://blog.dailydoseofds.com/p/10-regression-and-classification>

# Gradient Descent: How to Learn?



UNIVERSIDADE  
CATOLICA  
PORTUGUESA  
BRAGA



# Loss Optimization


- We want to find the network weights that achieve the lowest loss!

$$W^* = \operatorname{argmin}_W \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; W), y^{(i)})$$

$$W^* = \operatorname{argmin}_W J(W)$$

# Gradient Descent: Intuition

- Repeat:
  1. Guess
  2. Calculate error
- e.g., learn linear model for converting kilometers to miles when only observing the input “miles” and output “kilometers”



# Gradient Descent: Intuition

- Repeat:
  1. **Guess**
  2. Calculate error



A user interface for unit conversion. On the left, a box contains the number '1' above a dropdown menu labeled 'Mile'. To the right of this is an equals sign, followed by another box containing the number '1.60934' above a dropdown menu labeled 'Kilometer'. The box containing '1.60934' is highlighted with a red border.

1 → Kilometers = miles x **constant**



# Gradient Descent: Intuition

- Repeat:
  1. Guess
  2. Calculate error

1	=	1.60934
Mile		Kilometer



# Gradient Descent: Intuition

- Repeat:
  1. **Guess**
  2. Calculate error



A user interface for unit conversion. On the left, a box contains the number '1' above a dropdown menu labeled 'Mile'. To the right of this box is an equals sign. Further right is another box containing the number '1.60934' above a dropdown menu labeled 'Kilometer'. The box containing '1.60934' is highlighted with a thick red border.

1.5 → Kilometers = miles x constant

# Gradient Descent: Intuition

- Repeat:
  1. Guess
  2. Calculate error

1	=	1.60934
Mile		Kilometer

1.5 →  $\text{Kilometers} = \text{miles} \times \text{constant}$  →  $\text{Error} = \text{Guess} - \text{Correct}$

- Idea: iteratively adjust the **constant (i.e., model parameter)** to try to reduce the error

# Gradient Descent: Intuition


- Iteratively search for values that solve optimization problem (i.e., minimize or maximize an objective function)
- When **minimizing** the objective function, it also is often called interchangeably the **cost function**, **loss function**, or **error function**.

# Gradient Descent: Intuition

- Idea: use derivatives!
  - Derivatives tell us how a small change in the input  $x$  affects the output  $f(x)$ .
  - Specifically, the derivative  $f'(x)$  represents the rate of change of the function at a given point, indicating how much  $f(x)$  will change in response to a small change in  $x$ .
  - Gradient is a vector that indicates how  $f(x)$  changes as each function variable changes (i.e., partial derivatives).

# Gradient Descent: Intuition

- Gradient descent:
  - Iteratively take steps in the opposite direction of the gradient to minimize the function



Which letter(s) are the global minima?


Which letter(s) are local minima?



# Gradient Descent

$$W^* = \underset{W}{\operatorname{argmin}} J(W)$$

1. Randomly pick an initial  $(W_0, W_1)$



2. Compute the gradient:  $\frac{\partial J(W)}{\partial W}$

3. Take a small step in the opposite direction of the gradient.

4. Repeat until convergence.

# Gradient Descent


- Algorithm:
  1. Initialize weights randomly  $\sim N(0, \sigma^2)$
  2. Loop until convergence
  3. Compute gradient  $\frac{\partial J(W)}{\partial W}$
  4. Update weights  $W \leftarrow W - \eta \frac{\partial J(w)}{\partial w}$
  5. Return weights

# Computing Gradients: Backpropagation



- How does a small change in one weight (e.g.,  $w_2$ ) affect the final loss  $J(W)$ ?

# Computing Gradients: Backpropagation




$$\frac{\partial J(W)}{\partial w_2} =$$

Let's use the chain rule!

$$\frac{dy}{dx} = \frac{dy}{du} \cdot \frac{du}{dx}$$

$\frac{dy}{du}$ : Differentiate outer function, Keep the inside the same  
 $\frac{du}{dx}$ : Differentiate inner function


$$y(x) = (x^2 + 1)^3 \longrightarrow ?$$



$$\frac{\partial h}{\partial x_1} = \frac{\partial h}{\partial u_1} \cdot \frac{\partial u_1}{\partial x_1} + \frac{\partial h}{\partial u_2} \cdot \frac{\partial u_2}{\partial x_1}$$

$$\frac{\partial h}{\partial x_2} = \frac{\partial h}{\partial u_1} \cdot \frac{\partial u_1}{\partial x_2} + \frac{\partial h}{\partial u_2} \cdot \frac{\partial u_2}{\partial x_2}$$

# Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_2} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{purple bar}} * \underbrace{\frac{\partial \hat{y}}{\partial w_2}}_{\text{red bar}}$$

# Computing Gradients: Backpropagation




$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{red box}} * \underbrace{\frac{\partial \hat{y}}{\partial w_1}}_{\text{yellow box}}$$

↑  
Apply the chain rule!

↑  
Apply the chain rule!




# Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{blue}} * \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{yellow}} * \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{red}}$$

# Computing Gradients: Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{blue}} * \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{yellow}} * \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{red}}$$



Repeat this for **every weight in the network** using gradients from later layers.

# Gradient Descent: How Often to Update?

- **Batch Gradient Descent:**
  - Uses all training examples in each update
  - Less noisy updates but slow or infeasible for large datasets
- **Stochastic Gradient Descent (SGD):**
  - Updates with a single example per iteration
  - Fast and memory-efficient for large datasets, but updates can be noisy
- **Mini-batch Gradient Descent:**
  - Updates with a subset of examples per iteration
  - Reduces noise compared to SGD, scalable to large datasets, but can still be slow with large data

# Gradient Descent: How Much to Update?

- Step size = learning rate
  - (a) When learning rate is too small, convergence to good solution will be slow!
  - (b) When learning rate is too large, convergence to a good solution is not possible!



# Gradient Descent: How to Choose the Learning Rate?



UNIVERSIDADE  
CATOLICA  
PORTUGUESA  
BRAGA

## Idea 1

Try lots of different learning rates and see what works better!

# Gradient Descent: How to Choose the Learning Rate?



UNIVERSIDADE  
CATOLICA  
PORTUGUESA  
BRAGA

## Idea 3

Do something smarter!

Design an adaptive learning rate that "adapts" to the landscape!





# Adaptive Learning Rates

- Learning rates are no longer fixed!
- Can be made larger or smaller depending on:
  - How large the gradient is.
  - How fast learning is happening.
  - The size of particular weights.
  - etc...