



UNIVERSIDADE  
CATÓLICA  
PORTUGUESA

BRAGA

# Deep Learning

Session 9

## Regularization

Applied Data Science

2024/2025

# Regularization

- What is it?
  - A technique that constrains our optimization problem to discourage complex models by limiting the model's capacity, preventing it from fitting the noise in the training data.

“any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.”

- Ch. 5.2 of Goodfellow book on Deep Learning

- Why do we need it?
  - To improve the generalization of our model on unseen data by balancing model capacity, ensuring it captures the underlying patterns without overfitting.

# Preventing Overfitting

- **Get more data:** Increasing the amount of data helps the model generalize better.
- **Use a model with the right capacity:**
  - Too much capacity: The model becomes overly complex and learns noise (overfitting).
  - Too little capacity: The model fails to capture important patterns (underfitting).
- **Early stopping:** Stop training when the model starts to overfit, based on validation performance.
- **Parameter Norm Penalty:** Add a penalty to the loss function for large weights, discouraging complexity.

# Preventing Overfitting


- **Dropout:** Randomly drop neurons during training to prevent co-adaptation and reduce overfitting.
- **Batch Normalization:** Normalize activations in each layer, stabilizing learning and acting as a regularizer.
- **Ensemble methods:** Combine predictions from multiple models (like bagging or boosting) to reduce variance and improve generalization.

# Get More Data

- If possible, **gathering more data** is always the best solution!
  - However, data collection can be expensive or time-consuming.
  - More data may require more computational resources.
- **Data Augmentation:**
  - Data augmentation involves applying various transformations to your existing dataset to artificially increase its size and diversity.
- **Leverage Pre-trained Models:**
  - Transfer learning from models trained on larger datasets can help when gathering more data is difficult.


# Data Augmentation

**SMOTE**




**Random**

**Borderline-SMOTE**



**Borderline**


**ADASYN**



**Density**

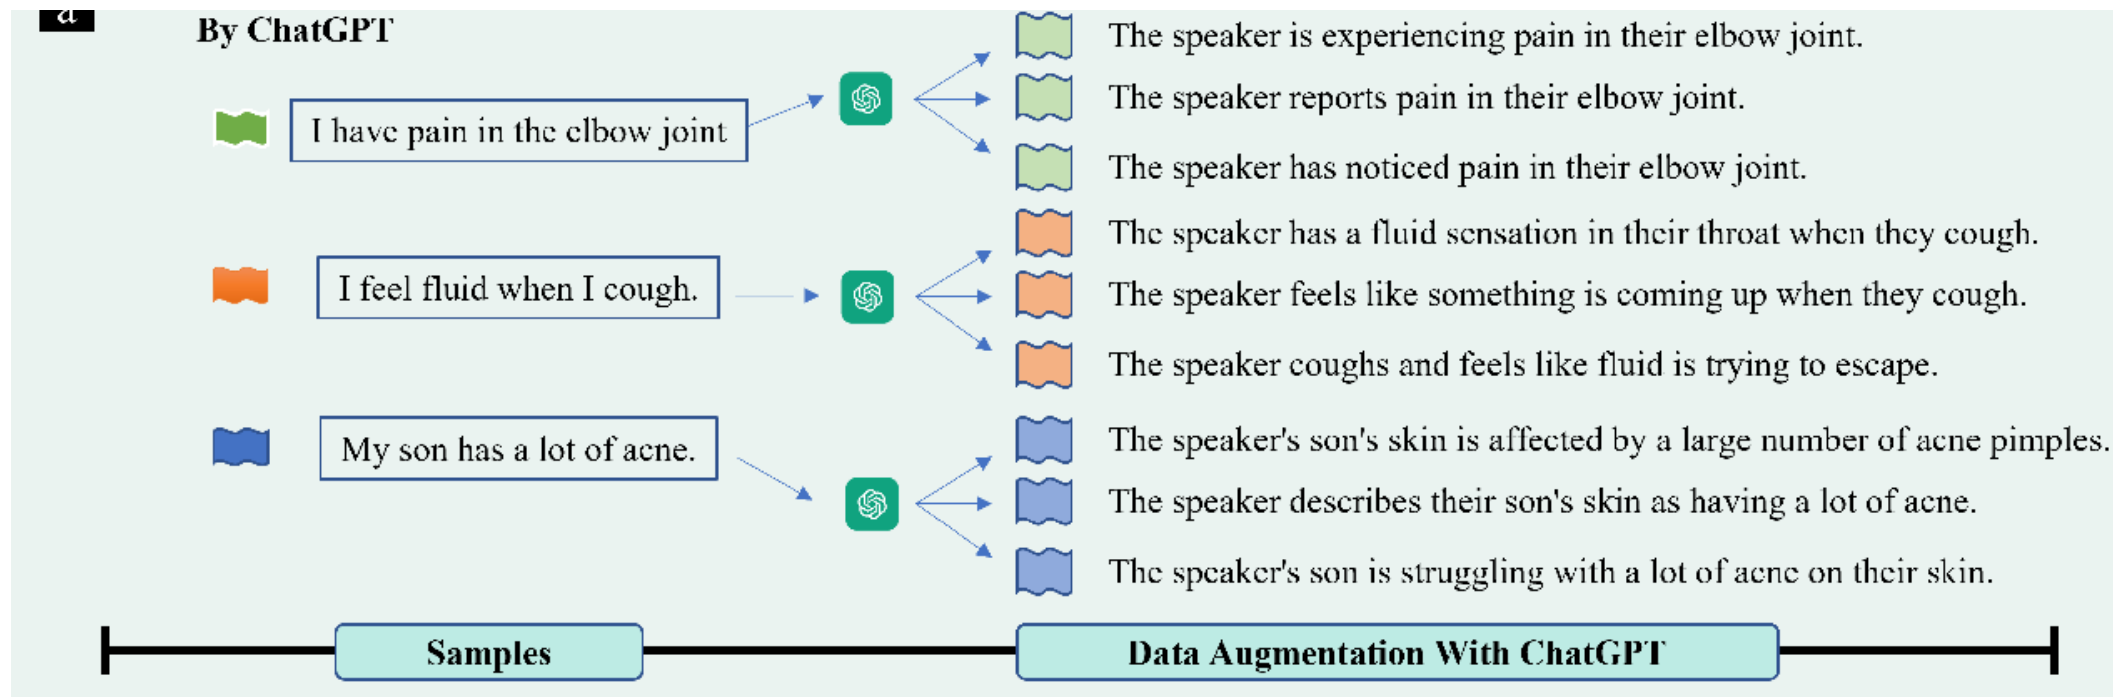
<https://towardsdatascience.com/smote-synthetic-data-augmentation-for-tabular-data-1ce28090debc>

# Data Augmentation



<https://www.labellerr.com/blog/what-is-data-augmentation-techniques-examples-benefits/>

# Data Augmentation





<https://www.catalyzex.com/paper/auggpt-leveraging-chatgpt-for-text-data>



# Early Stopping

- Stop training before we have a chance to overfit



# Parameter Norm Penalty

- **Key Insight:** Large weights are often a sign of overfitting.
- **Solution:** Apply a penalty to the size of the weights to reduce overfitting.
- **Analogy:** It's like tightening a belt on oversized pants.

# Parameter Norm Penalty

- **Idea:** Penalize large weights in the objective function
- e.g., objective is to minimize **sum of squared errors** over training examples
- **L2 norm (Ridge):** penalize squared weight values

$$Error = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \alpha \sum_{j=1}^m w_j^2$$

- **L1 norm (Lasso):** penalize absolute weight values

$$Error = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \alpha \sum_{j=1}^m |w_j|$$

- Note: only weights are penalized, not bias terms

# Parameter Norm Penalty

- e.g., objective is to minimize sum of squared errors over training examples.

- **L2 norm (Ridge):** penalize squared weight values

$$Error = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \boxed{\alpha} \sum_{j=1}^m w_j^2$$


- **L1 norm (Lasso):** penalize absolute weight values

$$Error = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 + \boxed{\alpha} \sum_{j=1}^m |w_j|$$

- **Hyperparameter** determines relative contribution of norm penalty term.

# Parameter Norm Penalty: How to Set Alpha?


- Shown is the same neural network with different levels of regularization. Which model has the largest value for alpha (i.e., largest norm penalty contribution)?




# Dropout

- Dropout is a stochastic regularization method.
- In each forward pass, randomly set some neurons to zero for one pass.
- Forces the network to not rely on any single node.
- The probability of dropping a neuron is defined by an hyperparameter; 0.5 is commonly used.
- Note: During inference dropout is not applied!

# Dropout



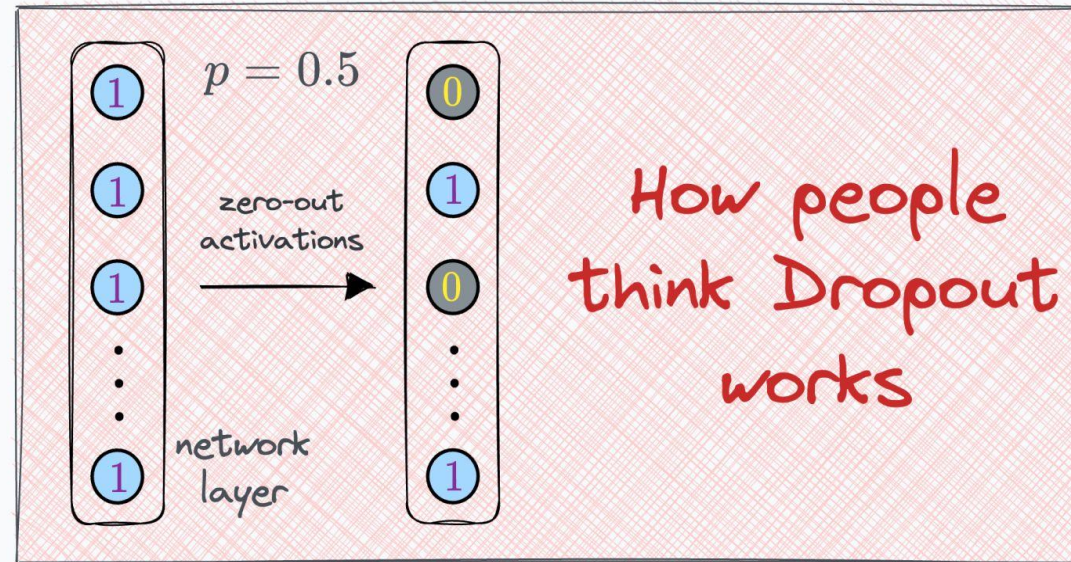
(a) Standard Neural Net




(b) After applying dropout.



# Dropout



 [blog.DailyDoseofDS.com](http://blog.DailyDoseofDS.com)






# Batch Normalization

- **Motivation:** Features on different scales can cause learning to be slower and poor performance
- We normalize all training data so that it resembles a normal distribution (that means, zero mean and a unitary variance)
- In the intermediate layers the distribution of the activations is constantly changing during training
  - This slows down the training process because each layer must learn to adapt themselves to a new distribution in every training step.
  - Batch normalization is a method we can use to normalize the inputs of each layer, in order to fight the internal covariate shift problem.

# Batch Normalization

- During training time, a batch normalization layer does the following:
  - Calculate the mean and variance of the layers input
  - Normalize the layer inputs using the previously calculated batch statistics
  - Scale and shift in order to obtain the output of the layer
  - $\gamma$  and  $\beta$  are learned during training along with the original parameters of the network.
- During inference, the mean and the variance are fixed. They are estimated using the previously calculated means and variances of each training batch.



**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;  
Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

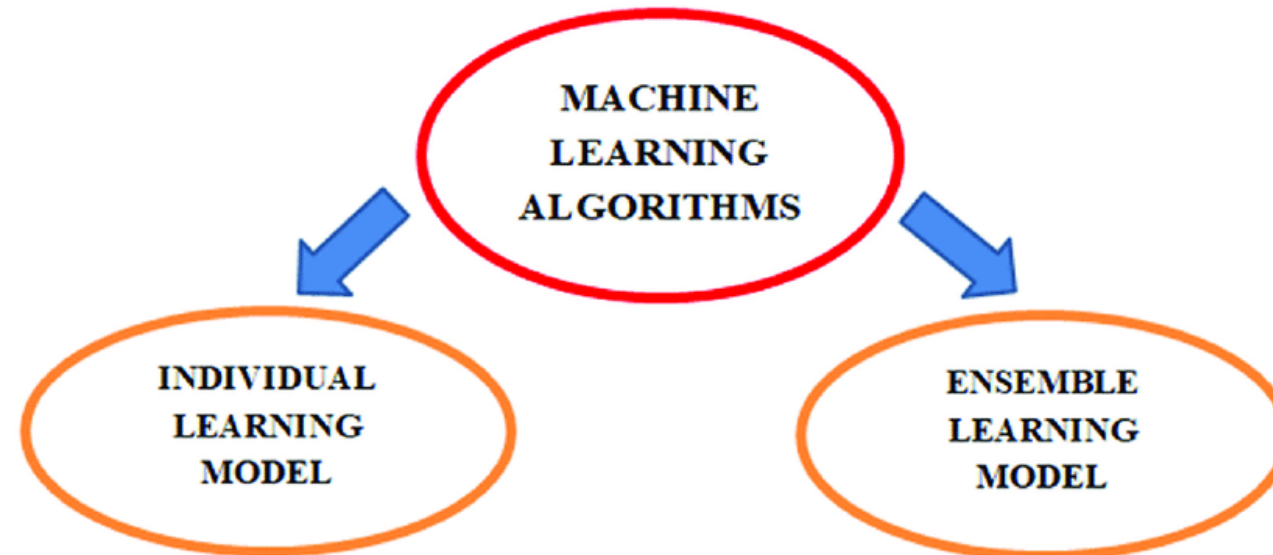
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

# Ensemble Methods

- **Idea:** Use the wisdom of the crowd
- **Why Choose Ensemble vs One Predictor?**
  - Reduces probability for making a wrong prediction



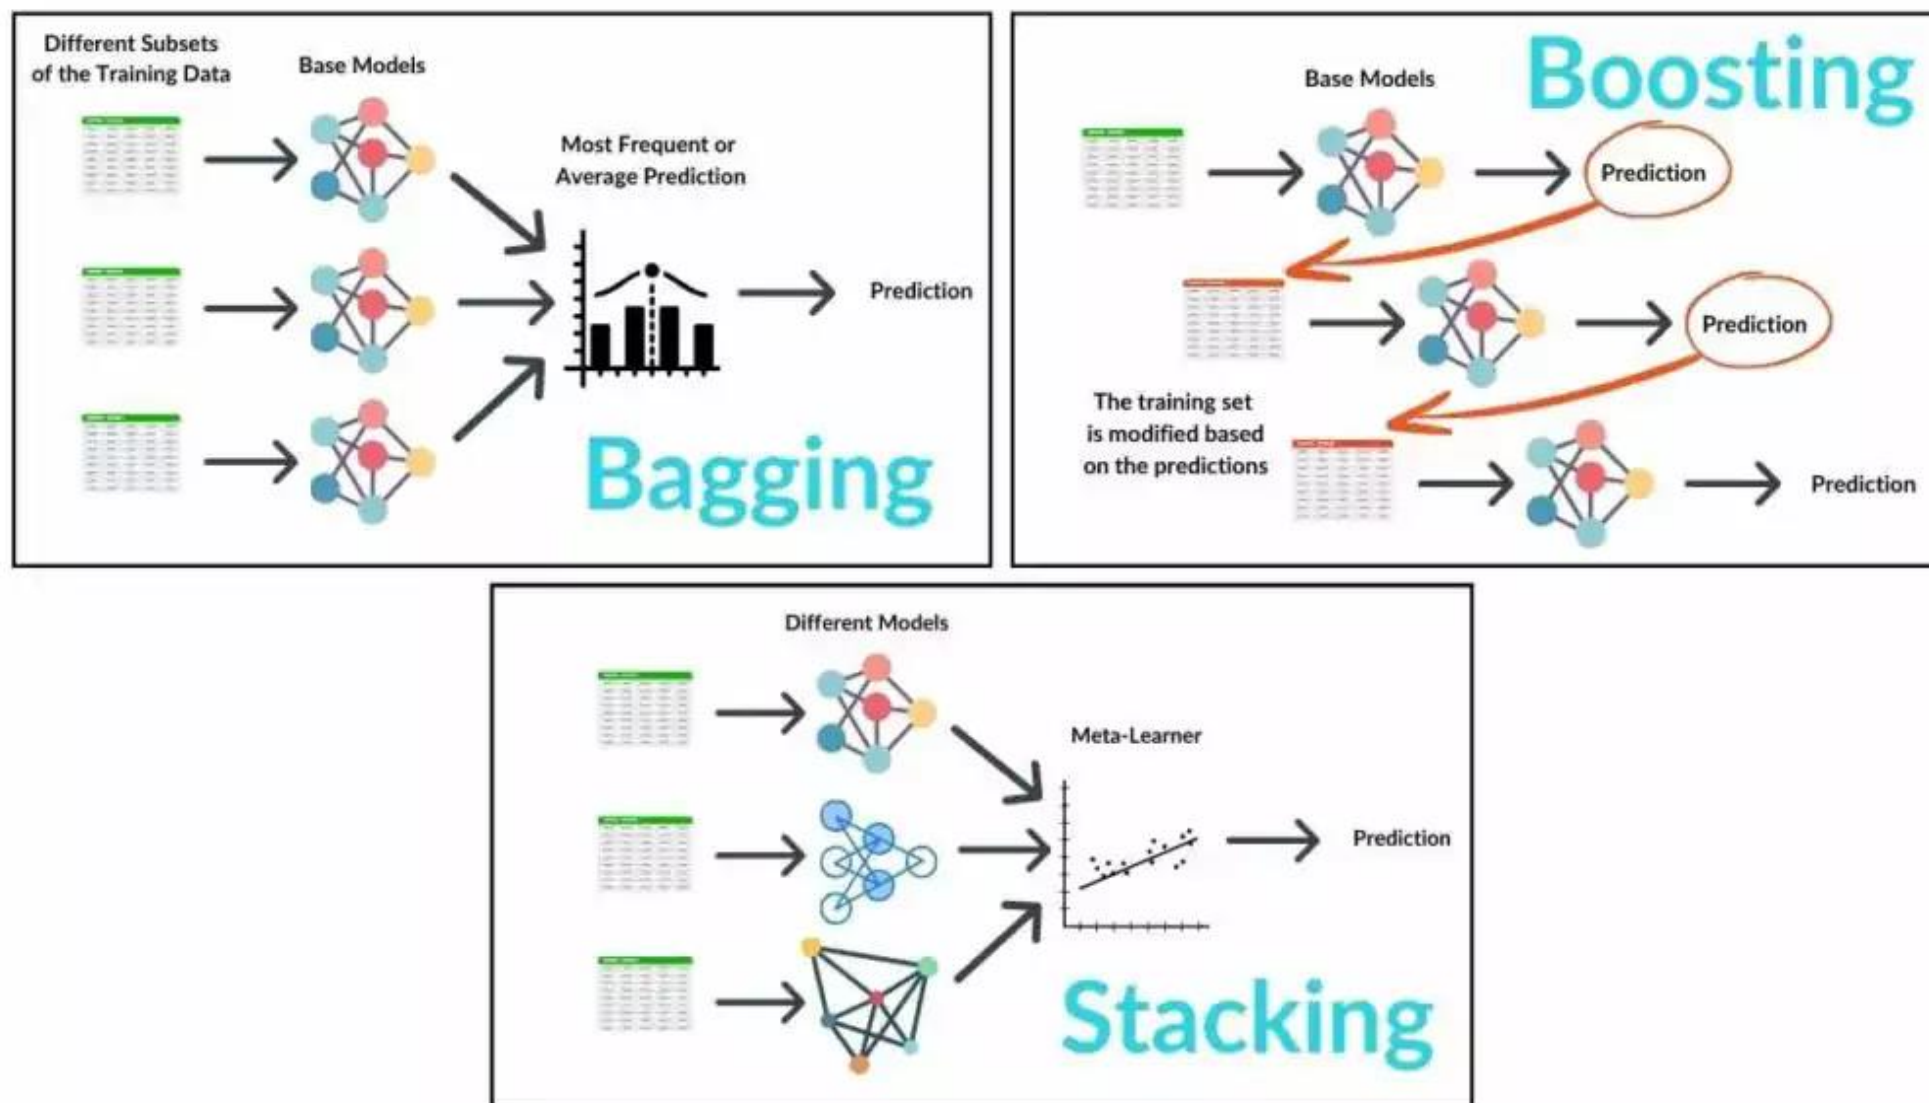
# Ensemble Methods

- Suppose:
  - n classifiers for binary classification task
  - Each classifier has same error rate  $\epsilon$
  - Classifiers are independent (not true in practice!)
  - Probability mass function indicates the probability of error from an ensemble:

$$P(y \geq k) = \sum_k^{\overbrace{n}^{\text{Number of classifiers}}} \underbrace{\binom{n}{k}}_{\substack{\text{\# ways to choose } k \text{ subsets from set of size } n}} \epsilon^k (1 - \underbrace{\epsilon}_{\text{Classifier error rate}})^{n-k} = \underbrace{\epsilon_{\text{ensemble}}}_{\text{Error probability}}$$

- e.g.,  $n = 11$ ,  $\epsilon = 0.25$ ;  $k = 6$ : probability of error is  $\sim 0.034$  which is much lower than probability of error from a single algorithm (0.25)

# How to Produce an Ensemble?




# Ensembles for Neural Networks

- Why could ensembling neural networks be difficult?
  - **Hyperparameter Tuning:** Finding optimal hyperparameters for each model is time-consuming.
  - **High Resource Usage:** Ensembles require significant memory and computational power.
  - **Extended Training Time:** Training multiple models increases total time.
  - **Increased Complexity:** Managing multiple architectures and data pipelines adds complexity.
  - **Diminishing Returns:** Additional models may only slightly improve performance.
  - **Deployment Issues:** Ensemble models increase inference time and hardware needs.




# Ensembles for Neural Networks

- **Idea:** approximate bagging with **dropout** during training so **different sub-models** in the network are trained with **different training data**



(a) Standard Neural Net



(b) After applying dropout.

e.g., drop 50% of  
units in hidden layers

e.g., drop 20% of  
units in input layers

# Regularization with Pytorch

- Dropout: <https://pytorch.org/docs/stable/generated/torch.nn.Dropout.html>
- L2 Regularization: <https://discuss.pytorch.org/t/how-to-add-a-l2-regularization-term-in-my-loss-function/17411/5>
- Batch Normalization: <https://medium.com/@aidant0001/batch-normalization-with-pytorch-959744b05325>
- Early Stopping: <https://www.geeksforgeeks.org/how-to-handle-overfitting-in-pytorch-models-using-early-stopping/#step-6-train-the-model-with-early-stopping>